

# untitled40

August 7, 2023

```
[1]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report
from sklearn.preprocessing import StandardScaler, LabelEncoder
```

```
[2]: df = pd.read_csv("C:/Users/DELL/Downloads/Dataset.csv")
df.head()
```

```
[2]:   encounter_id  patient_id  hospital_id  hospital_death    age    bmi \
0          66154      25312        118                  0  68.0  22.73
1         114252      59342        81                  0  77.0  27.42
2         119783      50777        118                  0  25.0  31.95
3          79267      46918        118                  0  81.0  22.64
4          92056      34377        33                  0  19.0     NaN

  elective_surgery  ethnicity gender  height ... aids cirrhosis \
0            0  Caucasian     M  180.3 ...  0.0      0.0
1            0  Caucasian     F  160.0 ...  0.0      0.0
2            0  Caucasian     F  172.7 ...  0.0      0.0
3            1  Caucasian     F  165.1 ...  0.0      0.0
4            0  Caucasian     M  188.0 ...  0.0      0.0

  diabetes_mellitus  hepatic_failure  immunosuppression  leukemia  lymphoma \
0           1.0            0.0                  0.0      0.0      0.0
1           1.0            0.0                  0.0      0.0      0.0
2           0.0            0.0                  0.0      0.0      0.0
3           0.0            0.0                  0.0      0.0      0.0
4           0.0            0.0                  0.0      0.0      0.0

  solid_tumor_with_metastasis  apache_3j_bodysystem  apache_2_bodysystem
0                      0.0                  Sepsis          Cardiovascular
1                      0.0                Respiratory          Respiratory
2                      0.0                Metabolic          Metabolic
3                      0.0  Cardiovascular          Cardiovascular
```

```
4
```

```
0.0
```

```
Trauma
```

```
Trauma
```

```
[5 rows x 186 columns]
```

```
[3]: df.dtypes
```

```
[3]: encounter_id           int64
patient_id            int64
hospital_id            int64
hospital_death         int64
age                   float64
...
leukemia              float64
lymphoma              float64
solid_tumor_with_metastasis float64
apache_3j_bodysystem   object
apache_2_bodysystem    object
Length: 186, dtype: object
```

```
[4]: df.columns
```

```
[4]: Index(['encounter_id', 'patient_id', 'hospital_id', 'hospital_death', 'age',
       'bmi', 'elective_surgery', 'ethnicity', 'gender', 'height',
       ...
       'aids', 'cirrhosis', 'diabetes_mellitus', 'hepatic_failure',
       'immunosuppression', 'leukemia', 'lymphoma',
       'solid_tumor_with_metastasis', 'apache_3j_bodysystem',
       'apache_2_bodysystem'],
      dtype='object', length=186)
```

```
[5]: # Handle Missing Values
# For numerical features, fill missing values with the mean
numeric_features = ['age', 'bmi']
df[numeric_features] = df[numeric_features].fillna(df[numeric_features].mean())
```

```
[6]: # Encode Categorical Features
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report
from sklearn.preprocessing import StandardScaler, LabelEncoder

encoder = LabelEncoder()
categorical_features = ['ethnicity', 'gender']
for col in categorical_features:
    df[col] = encoder.fit_transform(df[col].astype(str))
```

```
[23]: # Interactive Visualizations using Plotly Express

# Scatter plot to explore relationships between age and BMI
scatter_age_bmi = px.scatter(df, x='age', y='bmi', color='sepsis_label',
                               title='Relationship between Age, BMI, and Sepsis')
scatter_age_bmi.show()
```

  

```
[25]: # Parallel coordinates plot to visualize multiple features at once
parallel_plot = px.parallel_coordinates(df, dimensions=['age', 'bmi', 'gender'],
                                         color='sepsis_label',
                                         labels={'age': 'Age', 'bmi': 'BMI',
                                                 'ethnicity': 'Ethnicity', 'gender': 'Gender'},
                                         title='Parallel Coordinates Plot of Features and Sepsis')
parallel_plot.show()
```

  

```
[27]: # Box plot to compare the distribution of age among sepsis and non-sepsis cases
box_age = px.box(df, x='sepsis_label', y='age', title='Age Distribution among Sepsis and Non-Sepsis Cases')
box_age.show()
```

  

```
[28]: # Histogram to explore the distribution of BMI among sepsis and non-sepsis cases
hist_bmi = px.histogram(df, x='bmi', color='sepsis_label', title='BMI Distribution among Sepsis and Non-Sepsis Cases')
hist_bmi.show()
```

  

```
[8]: # Define sepsis label based on diagnoses or other relevant indicators
df['sepsis_label'] = (df['aids'] == 1) | (df['cirrhosis'] == 1) | (df['diabetes_mellitus'] == 1)
```

  

```
[9]: # Select relevant features for sepsis detection
selected_features = ['age', 'bmi', 'sepsis_label'] # Update with the correct feature names
data_selected = df[selected_features]
```

  

```
[10]: # Splitting data into features (X) and target (y)
X = data_selected.drop(['sepsis_label'], axis=1)
y = data_selected['sepsis_label']
```

  

```
[11]: # Normalize Numerical Features
scaler = StandardScaler()
X = scaler.fit_transform(X)

# Splitting data into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
random_state=42)
```

```
[12]: # Train and Evaluate Random Forest Classifier  
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)  
rf_model.fit(X_train, y_train)  
rf_predictions = rf_model.predict(X_test)  
rf_accuracy = accuracy_score(y_test, rf_predictions)  
print("Random Forest Accuracy:", rf_accuracy)  
print("Random Forest Classification Report:\n", classification_report(y_test,  
rf_predictions))
```

Random Forest Accuracy: 0.6654854712969526

Random Forest Classification Report:

	precision	recall	f1-score	support
False	0.77	0.80	0.78	13987
True	0.27	0.24	0.26	4356
accuracy			0.67	18343
macro avg	0.52	0.52	0.52	18343
weighted avg	0.65	0.67	0.66	18343

```
[13]: # Train and Evaluate Support Vector Machine (SVM)  
svm_model = SVC(kernel='linear', random_state=42)  
svm_model.fit(X_train, y_train)  
svm_predictions = svm_model.predict(X_test)  
svm_accuracy = accuracy_score(y_test, svm_predictions)  
print("SVM Accuracy:", svm_accuracy)  
print("SVM Classification Report:\n", classification_report(y_test,  
svm_predictions))
```

SVM Accuracy: 0.7625252139780843

SVM Classification Report:

	precision	recall	f1-score	support
False	0.76	1.00	0.87	13987
True	0.00	0.00	0.00	4356
accuracy			0.76	18343
macro avg	0.38	0.50	0.43	18343
weighted avg	0.58	0.76	0.66	18343

C:\Users\DELL\anaconda3\lib\site-  
packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning:  
Precision and F-score are ill-defined and being set to 0.0 in labels with no

```

predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\DELL\anaconda3\lib\site-
packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\DELL\anaconda3\lib\site-
packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))

```

[14]: # Train and Evaluate K-Nearest Neighbors (KNN)

```

knn_model = KNeighborsClassifier(n_neighbors=5)
knn_model.fit(X_train, y_train)
knn_predictions = knn_model.predict(X_test)
knn_accuracy = accuracy_score(y_test, knn_predictions)
print("KNN Accuracy:", knn_accuracy)
print("KNN Classification Report:\n", classification_report(y_test, knn_predictions))

```

KNN Accuracy: 0.7220192989151175

KNN Classification Report:

	precision	recall	f1-score	support
False	0.77	0.90	0.83	13987
True	0.32	0.15	0.20	4356
accuracy			0.72	18343
macro avg	0.54	0.52	0.52	18343
weighted avg	0.66	0.72	0.68	18343

[21]: # Create a DataFrame to store model names and accuracies

```

import plotly.express as px
import numpy as np

models = ['Random Forest', 'SVM', 'KNN']
accuracies = [rf_accuracy, svm_accuracy, knn_accuracy]
std_devs = [np.std(rf_predictions), np.std(svm_predictions), np.
            std(knn_predictions)]
model_data = pd.DataFrame({'Model': models, 'Accuracy': accuracies, 'Std Dev': std_devs})

# Plot Advanced-Level Graph using Plotly
fig = px.bar(model_data, x='Model', y='Accuracy', error_y='Std Dev',

```

```

        labels={'Accuracy': 'Accuracy Score', 'Std Dev': 'Standard
↳Deviation'},
        title='Model Accuracy Comparison',
↳color_discrete_sequence=['royalblue'])
fig.update_layout(xaxis={'categoryorder': 'total descending'})
fig.show()

```

```
[35]: import time # For simulating new data arrival

# Online Learning Simulation

for _ in range(10): # Simulate 10 new data points
    # Simulate new data point
    new_data_point = {
        'age': np.random.randint(18, 90),
        'bmi': np.random.uniform(15, 40),
        'ethnicity': np.random.choice(data['ethnicity'].unique()),
        'gender': np.random.choice(data['gender'].unique()),
        'sepsis_label': np.random.choice([True, False])
    }

    # Update models with the new data point
    new_data_df = pd.DataFrame([new_data_point])
    X_new = scaler.transform(new_data_df[['age', 'bmi']])
    y_new = new_data_df['sepsis_label'].values

    # Ensure that both classes are represented in the data
    if len(np.unique(y_new)) == 2:
        # Update RandomForest model with warm start
        rf_model.n_estimators += 1
        rf_model.fit(X_new, y_new)

        # Update SVM and KNN models with the new data point
        svm_model.fit(X_new, y_new)
        knn_model.fit(X_new, y_new)

    # Print updated model accuracies
    print("Random Forest Accuracy:", rf_model.score(X_test, y_test))
    print("SVM Accuracy:", svm_model.score(X_test, y_test))
    print("KNN Accuracy:", knn_model.score(X_test, y_test))

    # Pause to simulate data arrival interval
    time.sleep(3) # Wait for 3 seconds before simulating the next data point
```

Random Forest Accuracy: 0.7625252139780843

SVM Accuracy: 0.7625252139780843

```
[41]: # Perform k-fold cross-validation
from sklearn.model_selection import cross_val_score

k = 5 # Number of folds
cv_scores = cross_val_score(rf_model, X, y, cv=k, scoring='accuracy')

# Print cross-validation scores for each fold
for fold, score in enumerate(cv_scores, start=1):
    print(f"Fold {fold}: Accuracy = {score:.4f}")

# Calculate and print the mean and standard deviation of cross-validation scores
mean_score = np.mean(cv_scores)
std_dev_score = np.std(cv_scores)
print(f"Mean Accuracy: {mean_score:.4f}")
print(f"Standard Deviation: {std_dev_score:.4f}")
```

Fold 1: Accuracy = 0.6731  
Fold 2: Accuracy = 0.6670  
Fold 3: Accuracy = 0.6691

```
Fold 4: Accuracy = 0.6671  
Fold 5: Accuracy = 0.6727  
Mean Accuracy: 0.6698  
Standard Deviation: 0.0026
```

```
[51]: # Perform k-fold cross-validation for the SVM model  
svm_model = SVC(kernel='linear', random_state=42)  
k = 5 # Number of folds  
svm_cv_scores = cross_val_score(svm_model, X, y, cv=k, scoring='accuracy')
```

```
[52]: # Perform k-fold cross-validation for the KNN model  
knn_model = KNeighborsClassifier(n_neighbors=5)  
knn_cv_scores = cross_val_score(knn_model, X, y, cv=k, scoring='accuracy')
```

```
[53]: # Calculate and print the mean and standard deviation of cross-validation scores for SVM  
svm_mean_score = np.mean(svm_cv_scores)  
svm_std_dev_score = np.std(svm_cv_scores)  
print("SVM:")  
print(f"Mean Accuracy: {svm_mean_score:.4f}")  
print(f"Standard Deviation: {svm_std_dev_score:.4f}")  
  
# Calculate and print the mean and standard deviation of cross-validation scores for KNN  
knn_mean_score = np.mean(knn_cv_scores)  
knn_std_dev_score = np.std(knn_cv_scores)  
print("KNN:")  
print(f"Mean Accuracy: {knn_mean_score:.4f}")  
print(f"Standard Deviation: {knn_std_dev_score:.4f}")
```

```
SVM:  
Mean Accuracy: 0.7645  
Standard Deviation: 0.0000  
KNN:  
Mean Accuracy: 0.7238  
Standard Deviation: 0.0025
```

```
[44]: # Initialize individual models  
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)  
svm_model = SVC(kernel='linear', probability=True, random_state=42)  
knn_model = KNeighborsClassifier(n_neighbors=5)
```

```
[47]: # Initialize the VotingClassifier with the individual models  
ensemble_model = VotingClassifier(  
    estimators=[('rf', rf_model), ('svm', svm_model), ('knn', knn_model)],  
    voting='soft' # Use 'soft' voting for probability-based predictions  
)
```

```
[48]: # Train the ensemble model
ensemble_model.fit(X_train, y_train)

[48]: VotingClassifier(estimators=[('rf', RandomForestClassifier(random_state=42)),
                                    ('svm',
                                     SVC(kernel='linear', probability=True,
                                         random_state=42)),
                                    ('knn', KNeighborsClassifier())],
                        voting='soft')
```

```
[49]: # Evaluate the ensemble model
ensemble_accuracy = ensemble_model.score(X_test, y_test)
print("Ensemble Model Accuracy:", ensemble_accuracy)
```

Ensemble Model Accuracy: 0.7397917461702012